

Student Name: Mohammad Mohammad Beigi - Pantea Amoie
Student ID: 99102189 - 400101656
Subject: ML Privacy



Introduction to Machine Learning - Dr. R. Amiri

Final Project

1 Machine Unlearning

2 Private Training Attack

2.1 Base Model Learning

Train: X=(60000, 28, 28), y=(60000,)
Test: X=(10000, 28, 28), y=(10000,)
Model: "sequential_2"

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 784)	0
dense_4 (Dense)	(None, 128)	100,480
dense_5 (Dense)	(None, 10)	1,290

Total params: 101,770 (397.54 KB)
Trainable params: 101,770 (397.54 KB)
Non-trainable params: 0 (0.00 B)

Epoch 1/5
94/94 ————— 2s 4ms/step - accuracy: 0.5073 - loss: 82.7176
Epoch 2/5
94/94 ————— 1s 3ms/step - accuracy: 0.7224 - loss: 13.4081
Epoch 3/5
94/94 ————— 0s 3ms/step - accuracy: 0.7444 - loss: 8.7707
Epoch 4/5
94/94 ————— 1s 3ms/step - accuracy: 0.7929 - loss: 5.0982
Epoch 5/5
94/94 ————— 1s 3ms/step - accuracy: 0.7845 - loss: 1.9695
63/63 ————— 0s 2ms/step - accuracy: 0.6379 - loss: 3.7948

Model 0 test accuracy: 0.6439999938011169

Figure 1:

2.2 Base Model Shadow Learning and attack

```

Model 0 test accuracy: 0.643999938011169
Epoch 1/5
94/94 ----- 1s 3ms/step - accuracy: 0.4784 - loss: 90.4529
Epoch 2/5
94/94 ----- 1s 3ms/step - accuracy: 0.7028 - loss: 16.0662
Epoch 3/5
94/94 ----- 1s 3ms/step - accuracy: 0.7385 - loss: 9.1251
Epoch 4/5
94/94 ----- 0s 3ms/step - accuracy: 0.7479 - loss: 7.1349
Epoch 5/5
94/94 ----- 0s 3ms/step - accuracy: 0.7733 - loss: 5.2279
63/63 ----- 0s 2ms/step - accuracy: 0.6955 - loss: 7.1460

Model 0 test accuracy: 0.6790000200271606
Epoch 1/5
94/94 ----- 1s 3ms/step - accuracy: 0.4726 - loss: 82.6930
Epoch 2/5
94/94 ----- 1s 3ms/step - accuracy: 0.6843 - loss: 16.4903
Epoch 3/5
94/94 ----- 0s 3ms/step - accuracy: 0.7010 - loss: 5.9780
Epoch 4/5
94/94 ----- 1s 5ms/step - accuracy: 0.7113 - loss: 1.6879
Epoch 5/5
94/94 ----- 1s 5ms/step - accuracy: 0.7634 - loss: 1.2455
63/63 ----- 0s 3ms/step - accuracy: 0.7023 - loss: 2.9090

Model 1 test accuracy: 0.7080000042915344
Training svm model : 0
[LibSVM]SVM model 0 score : 0.6102150537634409

```

Figure 2: Number of shadow models = 2

2.3 Private Model Learning

```

Train: X=(60000, 28, 28), y=(60000,)
Test: X=(10000, 28, 28), y=(10000,)
Model: "sequential_4"

```

Layer (type)	Output Shape	Param #
flatten_4 (Flatten)	(None, 784)	0
dense_8 (Dense)	(None, 128)	100480
dropout_6 (Dropout)	(None, 128)	0
dense_9 (Dense)	(None, 10)	1290

```

=====
Total params: 101770 (397.54 KB)
Trainable params: 101770 (397.54 KB)
Non-trainable params: 0 (0.00 Byte)
=====
Epoch 1/5
94/94 [=====] - 1s 4ms/step - loss: 24.4152 - accuracy: 0.3330
Epoch 2/5
94/94 [=====] - 0s 3ms/step - loss: 2.3659 - accuracy: 0.4293
Epoch 3/5
94/94 [=====] - 0s 3ms/step - loss: 1.9294 - accuracy: 0.4753
Epoch 4/5
94/94 [=====] - 0s 4ms/step - loss: 1.7721 - accuracy: 0.5003
Epoch 5/5
94/94 [=====] - 0s 3ms/step - loss: 1.7717 - accuracy: 0.4990
63/63 [=====] - 0s 2ms/step - loss: 1.8605 - accuracy: 0.5755

Model 0 test accuracy: 0.5755000114440918

```

Figure 3: Private Trained Model by adding DropOut

2.4 Private Model Shadow Learning and attack

```

Model 0 test accuracy: 0.5755000114440918
Epoch 1/5
94/94 [=====] - 1s 3ms/step - loss: 28.6290 - accuracy: 0.3900
Epoch 2/5
94/94 [=====] - 0s 3ms/step - loss: 2.2916 - accuracy: 0.3817
Epoch 3/5
94/94 [=====] - 0s 3ms/step - loss: 1.9308 - accuracy: 0.4120
Epoch 4/5
94/94 [=====] - 0s 3ms/step - loss: 1.7087 - accuracy: 0.4530
Epoch 5/5
94/94 [=====] - 0s 4ms/step - loss: 1.6937 - accuracy: 0.4740
63/63 [=====] - 0s 2ms/step - loss: 1.3593 - accuracy: 0.6210

Model 0 test accuracy: 0.6209999918937683
Epoch 1/5
94/94 [=====] - 1s 3ms/step - loss: 24.4149 - accuracy: 0.3767
Epoch 2/5
94/94 [=====] - 0s 3ms/step - loss: 2.3103 - accuracy: 0.3637
Epoch 3/5
94/94 [=====] - 0s 4ms/step - loss: 1.9929 - accuracy: 0.4233
Epoch 4/5
94/94 [=====] - 1s 6ms/step - loss: 1.8375 - accuracy: 0.4540
Epoch 5/5
94/94 [=====] - 0s 5ms/step - loss: 1.7843 - accuracy: 0.4907
63/63 [=====] - 0s 3ms/step - loss: 1.3897 - accuracy: 0.5945

Model 1 test accuracy: 0.5945000052452087
Training svm model : 0
[LibSVM]SVM model 0 score : 0.48118279569892475

```

Figure 4: Number of shadow models = 2

2.5 Base Model Increasing Shadow Models

Model: "sequential_5"

Layer (type)	Output Shape	Param #
flatten_5 (Flatten)	(None, 784)	0
dense_10 (Dense)	(None, 128)	100,480
dropout_1 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 10)	1,290

Total params: 101,770 (397.54 KB)
Trainable params: 101,770 (397.54 KB)
Non-trainable params: 0 (0.00 B)

Epoch 1/5
94/94 [=====] 2s 8ms/step - accuracy: 0.4059 - loss: 85.0269
Epoch 2/5
94/94 [=====] 4s 35ms/step - accuracy: 0.4181 - loss: 2.2720
Epoch 3/5
94/94 [=====] 2s 3ms/step - accuracy: 0.4721 - loss: 1.8457
Epoch 4/5
94/94 [=====] 1s 3ms/step - accuracy: 0.4717 - loss: 2.0073
Epoch 5/5
94/94 [=====] 1s 14ms/step - accuracy: 0.5214 - loss: 1.6442
63/63 [=====] 1s 8ms/step - accuracy: 0.6274 - loss: 1.5189

Model 0 test accuracy: 0.6119999885559082
Epoch 1/5
94/94 [=====] 2s 6ms/step - accuracy: 0.3803 - loss: 74.3437
Epoch 2/5
94/94 [=====] 2s 23ms/step - accuracy: 0.3819 - loss: 2.2994
Epoch 3/5
94/94 [=====] 0s 4ms/step - accuracy: 0.4140 - loss: 1.8617
Epoch 4/5
94/94 [=====] 1s 12ms/step - accuracy: 0.4580 - loss: 1.7917
Epoch 5/5
94/94 [=====] 1s 4ms/step - accuracy: 0.4706 - loss: 1.5689
63/63 [=====] 0s 2ms/step - accuracy: 0.6694 - loss: 1.3016

Model 0 test accuracy: 0.6539999842643738

Model 0 test accuracy: 0.6539999842643738
Epoch 1/5
94/94 [=====] 2s 10ms/step - accuracy: 0.3444 - loss: 78.2884
Epoch 2/5
94/94 [=====] 1s 6ms/step - accuracy: 0.3500 - loss: 2.5628
Epoch 3/5
94/94 [=====] 0s 4ms/step - accuracy: 0.3818 - loss: 2.1204
Epoch 4/5
94/94 [=====] 0s 3ms/step - accuracy: 0.4212 - loss: 1.8821
Epoch 5/5
94/94 [=====] 1s 13ms/step - accuracy: 0.4528 - loss: 1.8724
63/63 [=====] 1s 6ms/step - accuracy: 0.5746 - loss: 1.5327

Model 1 test accuracy: 0.5590000152587891
Epoch 1/5
94/94 [=====] 3s 16ms/step - accuracy: 0.3453 - loss: 68.9072
Epoch 2/5
94/94 [=====] 1s 6ms/step - accuracy: 0.4100 - loss: 2.5006
Epoch 3/5
94/94 [=====] 2s 5ms/step - accuracy: 0.4345 - loss: 2.1591
Epoch 4/5
94/94 [=====] 0s 3ms/step - accuracy: 0.4741 - loss: 1.7842
Epoch 5/5
94/94 [=====] 2s 21ms/step - accuracy: 0.4934 - loss: 1.6147
63/63 [=====] 0s 2ms/step - accuracy: 0.6377 - loss: 1.9835

Model 2 test accuracy: 0.6305000185966492
Epoch 1/5
94/94 [=====] 3s 20ms/step - accuracy: 0.3645 - loss: 72.3915
Epoch 2/5
94/94 [=====] 1s 6ms/step - accuracy: 0.4231 - loss: 2.5470
Epoch 3/5
94/94 [=====] 2s 12ms/step - accuracy: 0.4441 - loss: 2.0397
Epoch 4/5
94/94 [=====] 0s 4ms/step - accuracy: 0.4795 - loss: 1.7557
Epoch 5/5
94/94 [=====] 1s 4ms/step - accuracy: 0.4982 - loss: 1.8641
63/63 [=====] 1s 15ms/step - accuracy: 0.6661 - loss: 1.3550

Model 3 test accuracy: 0.6614999771118164
Training svm model : 0
[LibSVM]SVM model 0 score : 0.667632665465435

Figure 5: Number of shadow models = 4

2.6 Safe Model Increasing Shadow Models

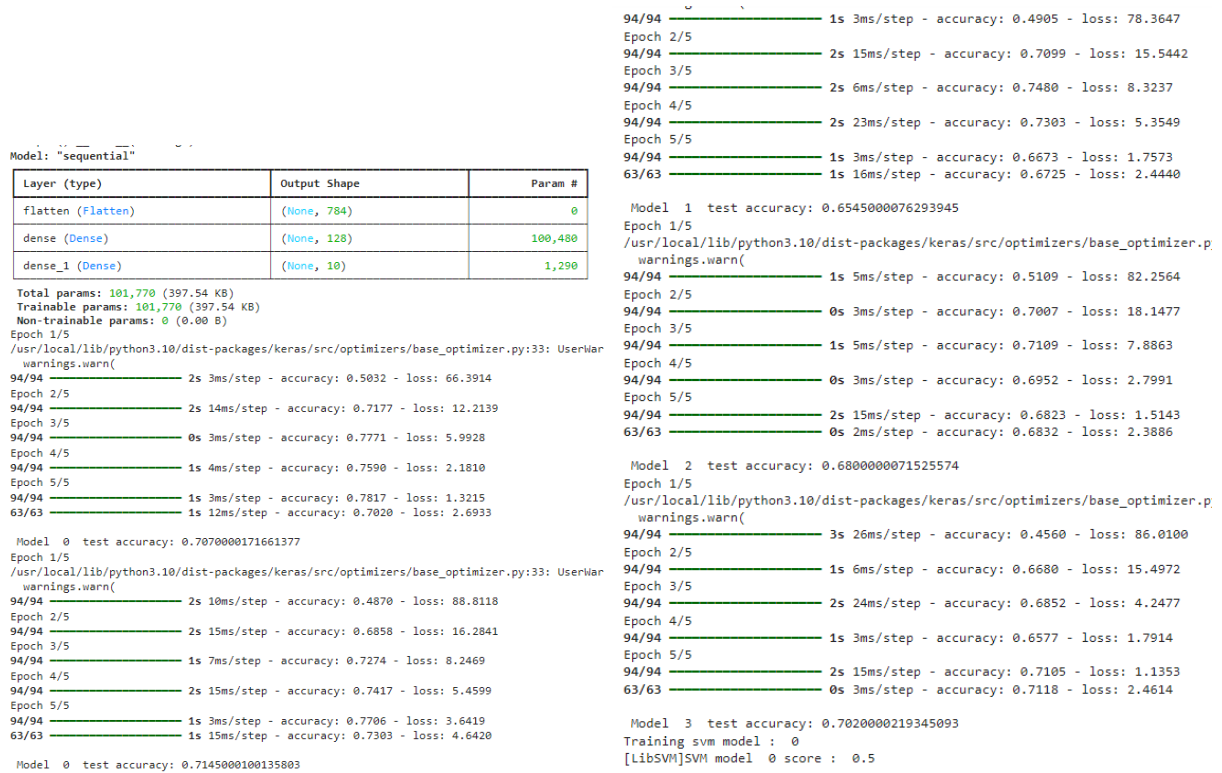


Figure 6: Number of shadow models = 4

3 Membership Inference Attack

We trained an attacker with 2 shadow models similar to Question 4, but using the new model.

```
1 from torchvision import models
2 import torch
3 import torchvision
4 import torchvision.transforms as transforms
5 from torch.utils.data import DataLoader
6 import torch.nn as nn
7 import torch.nn.functional as F
8 import torch.optim as optim
9 import numpy as np
10 import matplotlib.pyplot as plt
11
12
13 class CIFAR10Classifier(nn.Module):
14     def __init__(self):
15         super(CIFAR10Classifier, self).__init__()
16         self.conv1 = nn.Conv2d(3, 16, 3, 1)
17         self.conv2 = nn.Conv2d(16, 32, 3, 1)
18         self.dropout1 = nn.Dropout2d(0.25)
19         self.dropout2 = nn.Dropout2d(0.5)
20         self.fc1 = nn.Linear(6272, 64)
21         self.fc2 = nn.Linear(64, 10)
22
23     def forward(self, x):
24         x = self.conv1(x)
25         x = F.relu(x)
```

```

26 x = self.conv2(x)
27 x = F.relu(x)
28 x = F.max_pool2d(x, 2)
29 x = self.dropout1(x)
30 x = torch.flatten(x, 1)
31 x = self.fc1(x)
32 x = F.relu(x)
33 x = self.dropout2(x)
34 x = self.fc2(x)
35 return x
36
37 import tensorflow as tf
38 import tensorflow.keras
39 from tensorflow.keras.utils import to_categorical
40 import numpy as np
41 from sklearn.utils import resample
42 from sklearn.metrics import ConfusionMatrixDisplay
43 import matplotlib.pyplot as plt
44 from matplotlib import pyplot
45
46
47
48 def load_fashion_mnist_dataset():
49     from keras.datasets import fashion_mnist
50
51     # load dataset
52     (trainX, trainY), (testX, testY) = fashion_mnist.load_data()
53     # summarize loaded dataset
54     print('Train: X=%s, y=%s' % (trainX.shape, trainY.shape))
55     print('Test: X=%s, y=%s' % (testX.shape, testY.shape))
56
57     # reshape dataset to have a single channel
58     trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))
59     testX = testX.reshape((testX.shape[0], 28, 28, 1))
60
61     # one hot encode target values
62     trainY = to_categorical(trainY)
63     testY = to_categorical(testY)
64
65     return (trainX, trainY), (testX, testY)
66
67     # Set hyper-parameters for our training of neural networks
68     LEARNING_RATE = 0.001
69     EPOCH = 5
70     VERBOSE = 1
71
72     TRAINING_SIZE = 3000
73     TEST_SIZE = 2000
74
75     # No of target models
76     NUM_TARGET = 1
77     # No of shadow models
78     NUM_SHADOW = 2
79
80     # Label value "in" for records present in training data of shadow models
81     IN = 1
82     # Label value "out" for records not present in training data of shadow models
83     OUT = 0
84
85     def sample_data(train_data, test_data, num_sets):
86         (x_train, y_train), (x_test, y_test) = train_data, test_data

```

```

87 new_x_train, new_y_train = [], []
88 new_x_test, new_y_test = [], []
89 for i in range(num_sets):
90     x_temp, y_temp = resample(x_train, y_train, n_samples=TRAINING_SIZE, random_state=0)
91     new_x_train.append(x_temp)
92     new_y_train.append(y_temp)
93     x_temp, y_temp = resample(x_test, y_test, n_samples=TEST_SIZE, random_state=0)
94     new_x_test.append(x_temp)
95     new_y_test.append(y_temp)
96     return (new_x_train, new_y_train), (new_x_test, new_y_test)
97
98 def get_attack_dataset(models, train_data, test_data, num_models, data_size):
99     # generate dataset for the attack model
100     (x_train, y_train), (x_test, y_test) = train_data, test_data
101     # set number of classes for the attack model
102     num_classes = 10
103     x_data, y_data = [[] for _ in range(num_classes)], [[] for _ in range(num_classes)]
104     for i in range(num_models):
105         # IN data
106         x_temp, y_temp = resample(x_train[i], y_train[i], n_samples=data_size, random_state=0)
107         for j in range(data_size):
108             y_idx = np.argmax(y_temp[j])
109             x_data[y_idx].append(models[i].predict(x_temp[j:j+1], verbose=0)[0])
110             y_data[y_idx].append(IN)
111
112         # OUT data
113         x_temp, y_temp = resample(x_test[i], y_test[i], n_samples=data_size, random_state=0)
114         for j in range(data_size):
115             y_idx = np.argmax(y_temp[j])
116             x_data[y_idx].append(models[i].predict(x_temp[j:j+1], verbose=0)[0])
117             y_data[y_idx].append(OUT)
118
119     return x_data, y_data
120
121 def build_fcnn_model_fashion_mnist():
122     model = tf.keras.models.Sequential([
123         tf.keras.layers.Flatten(input_shape=(28, 28)),
124         tf.keras.layers.Dense(128, activation='relu'),
125         tf.keras.layers.Dense(10, activation='softmax')
126     ])
127     model.summary()
128     return model
129
130 def get_trained_keras_models(keras_model, train_data, test_data, num_models):
131     (x_train, y_train), (x_test, y_test) = train_data, test_data
132     models = []
133     for i in range(num_models):
134         models.append(tf.keras.models.clone_model(keras_model))
135         rms = tf.keras.optimizers.RMSprop(learning_rate=LEARNING_RATE, decay=1e-7)
136         models[i].compile(loss='categorical_crossentropy', optimizer=rms, metrics=['accuracy'])
137         models[i].fit(x_train[i], y_train[i], batch_size=32, epochs=EPOCH, verbose=VERBOSE,
138             shuffle=True)
139         score = models[i].evaluate(x_test[i], y_test[i], verbose=VERBOSE)
140         print('\n', 'Model ', i, ' test accuracy:', score[1])
141     return models
142
143 def get_trained_svm_models(train_data, test_data, num_models=1):
144     from sklearn import svm
145     (x_train, y_train), (x_test, y_test) = train_data, test_data
146     models = []
147     for i in range(num_models):

```

```

147 print('Training svm model : ', i)
148 models.append(svm.SVC(gamma='scale', kernel='linear', verbose=VERBOSE))
149 models[i].fit(x_train[i], y_train[i])
150 score = models[i].score(x_test[i], y_test[i])
151 print('SVM model ', i, 'score : ', score)
152 return models
153
154 def membership_attack():
155 # load the pre-shuffled train and test data
156 (x_train, y_train), (x_test, y_test) = load_fashion_mnist_dataset()
157
158 # split the data for each model
159 target_train = (x_train[:TRAINING_SIZE*NUM_TARGET], y_train[:TRAINING_SIZE*NUM_TARGET])
160 target_test = (x_test[:TEST_SIZE*NUM_TARGET], y_test[:TEST_SIZE*NUM_TARGET])
161 target_train_data, target_test_data = sample_data(target_train, target_test, NUM_TARGET)
162
163 shadow_train = (x_train[TRAINING_SIZE*NUM_TARGET:], y_train[TRAINING_SIZE*NUM_TARGET:])
164 shadow_test = (x_test[TEST_SIZE*NUM_TARGET:], y_test[TEST_SIZE*NUM_TARGET:])
165 shadow_train_data, shadow_test_data = sample_data(shadow_train, shadow_test, NUM_SHADOW)
166
167 cnn_model = build_fcnn_model_fashion_mnist()
168
169 # compile the target model
170 target_models = get_trained_keras_models(cnn_model, target_train_data, target_test_data,
171 NUM_TARGET)
172 # compile the shadow models
173 shadow_models = get_trained_keras_models(cnn_model, shadow_train_data, shadow_test_data,
174 NUM_SHADOW)
175
176 # get train data for the attack model
177 attack_train = get_attack_dataset(shadow_models, shadow_train_data, shadow_test_data,
178 NUM_SHADOW, TEST_SIZE)
179 # get test data for the attack model
180 attack_test = get_attack_dataset(target_models, target_train_data, target_test_data,
181 NUM_TARGET, TEST_SIZE)
182
183 # training the attack model
184 #attack_model = get_trained_svm_models(attack_train, attack_test)
185 return attack_train, attack_test
186 NUM_SHADOW = 2
187 membership_attack()
188
189
190 import torch
191 import torch.nn as nn
192 import torch.optim as optim
193 from torchvision.datasets import CIFAR10
194 from torchvision import transforms
195 from torch.utils.data import Subset, DataLoader, TensorDataset
196 from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
197 from sklearn.linear_model import LogisticRegression
198
199 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
200
201 model = CIFAR10Classifier()
202 state_dict = torch.load("model_state_dict.pth", map_location=device)
203 new_state_dict = {key.replace('_module.', ''): value for key, value in state_dict.items()}
204 model.load_state_dict(new_state_dict)
205 model.to(device)
206 model.eval()
207

```

```

204 transform = transforms.Compose([
205     transforms.ToTensor(),
206     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
207 ])
208
209 DATA_ROOT = '../cifar10'
210 BATCH_SIZE = 64
211
212 # Load the indices from list.txt
213 indices_file = 'list.txt' #####
214 with open(indices_file, 'r') as f:
215     indices = [int(line.strip()) for line in f]
216
217 full_train_dataset = CIFAR10(root=DATA_ROOT, train=True, download=True, transform=
    transform)
218 test_dataset = CIFAR10(root=DATA_ROOT, train=False, download=True, transform=transform)
219
220 train_indices_set = set(indices)
221 all_indices = set(range(len(full_train_dataset)))
222 other_indices = list(all_indices - train_indices_set)
223
224 train_dataset = Subset(full_train_dataset, indices[:len(indices)//2]) #####
225 other_dataset = Subset(full_train_dataset, other_indices)
226
227 # Create data loaders
228 train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=False)
229 other_loader = DataLoader(other_dataset, batch_size=BATCH_SIZE, shuffle=False)
230 test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)
231
232 # Create labels
233 train_labels = torch.ones(len(train_dataset)).to(device)
234 other_labels = torch.zeros(len(other_dataset)).to(device)
235 test_labels = torch.zeros(len(test_dataset)).to(device)
236 #####
237 #if you have an attacker model for each class, modify the above code.
238 #####
239
240 def extract_features(model, dataloader):
241     model.eval()
242     features = []
243     with torch.no_grad():
244         for data in dataloader:
245             inputs, _ = data
246             inputs = inputs.to(device)
247             outputs = model(inputs)
248             features.append(outputs)
249     return torch.cat(features).to(device)
250
251 train_features = extract_features(model, train_loader)
252 other_features = extract_features(model, other_loader)
253 test_features = extract_features(model, test_loader)
254
255
256 combined_features = torch.cat((train_features, other_features, test_features))
257 combined_labels = torch.cat((train_labels, other_labels, test_labels))
258
259
260 new_dataset = TensorDataset(combined_features, combined_labels)
261 new_loader = DataLoader(new_dataset, batch_size=BATCH_SIZE, shuffle=True)
262
263 #load your attacker model

```



```

264
265 #####
266
267 # Calculate training accuracy, confusion matrix, precision, and recall
268 binary_classifier.eval()
269 all_labels = []
270 all_predicted = []
271 correct = 0
272 total = 0
273 with torch.no_grad():
274     for features, labels in new_loader:
275         features, labels = features.to(device), labels.to(device)
276         outputs = get_trained_svm_models(features, attack_test).squeeze()
277         predicted = (outputs > 0.5).float()
278         total += labels.size(0)
279         correct += (predicted == labels).sum().item()
280         all_labels.extend(labels.cpu().numpy())
281         all_predicted.extend(predicted.cpu().numpy())
282
283 accuracy = correct / total
284
285
286 cm = confusion_matrix(all_labels, all_predicted)
287 precision = precision_score(all_labels, all_predicted)
288 recall = recall_score(all_labels, all_predicted)
289 f1 = f1_score(all_labels, all_predicted)
290
291 print(f'Confusion Matrix:\n{cm}')
292 print(f"Precision: {precision:.4f}")
293 print(f"Recall: {recall:.4f}")
294 print(f"F1 Score: {f1:.4f}")
295 print(f'Training Accuracy: {accuracy:.4f}')

```

```

Confusion Matrix:
[[ 9452  8766]
 [ 8361 13421]]
Precision: 0.6049
Recall: 0.6162
F1 Score: 0.6105
Training Accuracy: 0.5718

```

Figure 7:

PyTorch and CIFAR-10 Classifier

Explanation:

1. Imports:

- torchvision, torch, transforms: For handling data transformations and model-related functionalities.
- nn, F: For defining neural network layers and activation functions.
- optim: For optimization algorithms.
- numpy, matplotlib.pyplot: For numerical operations and plotting.

2. CIFAR10Classifier Class:

- Inherits from `nn.Module`.
- **Initialization** (`__init__`): Defines the structure of the neural network.
 - Convolutional layers (`conv1`, `conv2`), dropout layers (`dropout1`, `dropout2`), and fully connected layers (`fc1`, `fc2`).
- **Forward Method**: Defines the forward pass through the network.
 - Applies convolution, activation (ReLU), pooling, dropout, and flattening operations sequentially.

TensorFlow and CIFAR 10 Dataset

Explanation:

1. Imports:

- `tensorflow`, `keras.utils`: For building and managing neural networks.
- `numpy`, `sklearn.utils`, `sklearn.metrics`, `matplotlib.pyplot`: For data handling, resampling, and plotting.

2. `load_fashion_mnist_dataset` Function:

- Loads the Fashion-MNIST dataset.
- Reshapes the dataset to have a single channel.
- Converts the labels to categorical format (one-hot encoding).
- Returns the training and testing datasets.

Sampling and Training

Explanation:

1. `sample_data` Function:

- Resamples training and testing datasets to create multiple sets.
- Helps in creating diverse training and testing sets for training models.

2. `get_attack_dataset` Function:

- Generates datasets for training attack models.
- Uses predictions from multiple models on resampled data to create in-distribution (IN) and out-of-distribution (OUT) samples.

Neural Network and SVM Models

Explanation:

1. `build_fcnn_model_fashion_mnist` Function:

- Builds a fully connected neural network for Fashion-MNIST.
- Uses two dense layers with ReLU and softmax activations.

2. **get_trained_keras_models Function:**

- Trains multiple instances of a given Keras model.
- Uses RMSprop optimizer and evaluates each model.

3. **get_trained_svm_models Function:**

- Trains multiple SVM models using the provided training data.
- Evaluates each model and prints the accuracy.

Membership Attack

Explanation:

1. **membership_attack Function:**

- Implements a membership inference attack.
- Loads and preprocesses the Fashion-MNIST dataset.
- Splits data for training target and shadow models.
- Builds and trains neural networks (both target and shadow models).
- Creates datasets for training an attack model using predictions from shadow models.
- Returns the attack training and testing datasets.

Evaluating CIFAR-10 Classifier with Torch

Explanation:

1. **Device Setup:**

- Checks for GPU availability and sets the device accordingly.

2. **Model Loading:**

- Loads the pretrained CIFAR10Classifier model state.

3. **Data Preparation:**

- Defines transformations for the CIFAR-10 dataset.
- Loads CIFAR-10 dataset and splits it based on indices from a file.

4. **Feature Extraction:**

- Extracts features from the trained model for training, other, and test datasets.

5. **Creating New Dataset:**

- Combines extracted features and labels into a new dataset.
- Creates a DataLoader for the new dataset.

6. **Evaluation:**

- Evaluates the binary classifier on the new dataset.
- Computes accuracy, confusion matrix, precision, recall, and F1 score.